

Spatiotemporal Graph Neural Networks for Probabilistic Postprocessing of Wind Speed

Edoardo Leali Matteo Vitali Sidharth Padmanabhan


USI – MeteoSwiss

January 5, 2026

The Problem: NWP Systematic Biases

Challenges in NWP models:

- **Terrain resolution**
 - Grid smoothing flattens mountains
- **Precipitation biases**
 - "Drizzle" bias: too much light rain. Underestimate extreme events.
- **Temperature biases**
 - Damped day-night temperature range



Our Task

Statistical postprocessing to correct these systematic biases

Key Distinction

We're **NOT forecasting** from scratch → We're **correcting** existing forecasts

Dataset & Graph Structure

MeteoSwiss Wind Speed Dataset

- **Train/Val:** Reforecasts (Feb 2020 – Sep 2023)
- **Test:** Operational forecasts (May 2024 – Jan 2025)
- 18 predictors
- 96-hour forecast horizon

Spatial Graph Construction:

- **Nodes:** Weather stations
- **Edges:** k-NN ($k=5$)
 - Haversine distance
 - Gaussian kernel weights
- **Fixed graph**



Graph with the weather stations

The Machine Learning Task

Problem Formulation:

Input: $\mathbf{X} \in \mathbb{R}^{T \times N \times F}$ spatiotemporal data

- $T = 96$ timesteps (lead times 1–96 hours)
- $N =$ number of stations (nodes in graph)
- $F = 18$ features:
 - 4 NWP ensemble features (mean, std, pressure gradients)
 - 4 temporal features (hour-of-day, day-of-year encodings)
 - 10 terrain features (elevation, roughness, etc.)

Output: Probabilistic distribution

$$y_{s,t} \sim \text{LogNormal}(\mu_{s,t}, \sigma_{s,t})$$

- Why LogNormal? Wind speeds are strictly positive and right-skewed

Learning Objective

Supervised Learning Framework:

- **Input:** $\mathbf{X} \in \mathbb{R}^{T \times N \times F}$ (predictor features)
- **Output:** $(\mu_{s,t}, \sigma_{s,t})$ parameters for LogNormal distribution
- **Target:** $y_{s,t}$ (observed wind speed at station s , lead time t)
- **Loss:** CRPS (Continuous Ranked Probability Score)

CRPS Properties:

- Proper scoring rule: jointly optimizes sharpness & calibration
- Masked implementation: excludes missing observations (sensor failures)

$$\text{CRPS}(F, y) = \int_{-\infty}^{\infty} [F(t) - H(t - y)]^2 dt$$

where F is the predictive CDF, y is the observation, and $H(t - y)$ is the Heaviside function (0 when $t < y$, 1 otherwise)

Complete Experimental Pipeline

1. Model Implementation

- ① **Model0:** Pure temporal (no graph) – *baseline*
- ② **STGNN1:** Bidirectional RNN where the recurrent cell is a MP cell
- ③ **STGNN2:** TCN-GNN
- ④ **STGNN3:** Temporal Graph MLP-Mixer (our model)

2. Hyperparameter Tuning

- Optuna framework with TPE sampling
- Search space: learning rate, hidden channels, layers, dropout

3. Multi-seed Evaluation

- 5 independent random seeds for statistical robustness
- Report mean \pm std across seeds

4. Calibration Analysis

- Rank histograms (Talagrand diagrams)
- Assess predictive distribution calibration

Metrics: CRPS (probabilistic) + MAE (deterministic)

Model Implementations

Baseline - Enhanced Architectures - Our Model

Recap: Temporal Convolutional Networks

A Temporal Convolutional Network processes sequential data using **dilated 1D convolutions** instead of recurrent units.

Key Innovation:

- **Exponential dilation:** $d = 2^l$ for layer l
 - Layer 1: dilation = 1 (hourly patterns)
 - Layer 2: dilation = 2
 - Layer 3: dilation = 4
 - Layer 6: dilation = 32 (multi-day)

Advantages:

- Parallelizable (unlike RNNs)
- No vanishing gradients
- Multi-scale temporal patterns
- For postprocessing: non-causal convolutions
 - Look forward & backward in forecast window

Residual Connections:

- output = Layer(input) + input
- Network learns *corrections*, not full transformations
- Provides gradient highway → stable deep training

Model0: Non-Graph Temporal Baseline

Architecture: Temporal Convolutional Network (No Graph)

Forward Pass:

- ① **Encoder:** Linear projection $\mathbb{R}^{18} \rightarrow \mathbb{R}^{128}$
 - Projects input features to hidden dimension
- ② **Rearrange for TCN:** $[B, T, N, C] \rightarrow [(B \times N), C, T]$
 - Treats each station as **independent sequence**
 - N stations processed in parallel, no spatial info shared
- ③ **Stacked TCN Layers:** 2 layers with dilations [1, 2]
 - Each layer returns: $(x_{\text{residual}}, x_{\text{skip}})$
 - Collects multi-scale temporal features
- ④ **Multi-Scale Aggregation:** $\sum_{I=1}^L \text{skip}_I + x_{\text{final}}$
 - Combines all temporal scales into final representation
- ⑤ **Rearrange Back:** $[(B \times N), C, T] \rightarrow [B, T, N, C]$
- ⑥ **Probabilistic Output:** Projects to $\text{LogNormal}(\mu, \sigma)$ parameters

Model0: Performance

Results:

- **CRPS: 0.67** vs. NWP 0.98
- Strong improvement over raw ensemble
- *Most gain comes from learning temporal correction?*

Key Insight

This purely temporal baseline is **surprisingly good!**

→ Graph structure needs to prove its worth

Challenge for graph models: Beat this strong baseline by capturing spatial dependencies

STGNN1: What the Baseline Does

Architecture: Bidirectional Recurrent GNN

Approach:

- ① At each timestep t : aggregate spatial neighbors
- ② Then: update temporal state with RNN
- ③ Process both directions:
 - Forward: $t = 1 \rightarrow 96$
 - Backward: $t = 96 \rightarrow 1$

Components:

- Simple message passing (GatedGraphNetwork)
- Uniform neighbor weighting
- Additive residuals: $\mathbf{h}_t = \mathbf{h}_{t-1} + \Delta\mathbf{h}_t$

STGNN1: Problems & Solutions

Limitations:

- **Vanishing Gradients**

- Additive residuals:

$$\mathbf{h}_t = \mathbf{h}_{t-1} + \Delta\mathbf{h}_t$$

- 96 timesteps → exponential decay

- **Equal Neighbor Weighting**

- All neighbors treated equally
- Wind is directional – upwind matters more!

- **No Hierarchical Learning**

- Flat spatial processing
- Early vs. deep layers should learn different scales

Our Improvements:

- **GRU Cells**

- Gated updates:

$$\mathbf{h}_t = (1 - z_t) \odot \mathbf{h}_{t-1} + z_t \odot \tilde{\mathbf{h}}_t$$

- Robust gradient flow

- **GATv2 Attention**

- Dynamic neighbor weighting
- 2 attention heads

- **Layer-by-Layer Processing**

- Deep spatial learning
- Hierarchical pattern extraction

STGNN1: The Gradient Problem & Solution

Before:

Baseline used additive residuals:

$$\mathbf{h}_t = \mathbf{h}_{t-1} + \Delta\mathbf{h}_t$$

Result: Exponential gradient decay during backpropagation → model cannot learn.

After:

Replace with GRU cells:

$$\mathbf{h}_t = (1 - z_t) \odot \mathbf{h}_{t-1} + z_t \odot \tilde{\mathbf{h}}_t$$

Gating creates **multiplicative skip connections** → gradients flow without decay. Model finally learns temporal patterns




Figure 2. BiDirectional STGNN (STGNN1 default) at epoch 50.




Figure 3. Our tuned STGNN1 at epoch 50.

STGNN1: Weather Perspective

How Our Improvements Address Weather Postprocessing:

- **GRU Cells**

- Track sustained weather events (Föhn episodes, persistent winds)
- Maintain temporal context across 96-hour forecast window

- **GATv2 Attention**

- Identifies relevant upwind stations for wind correction: upwind stations weighted more heavily
- Captures local effects: valley channeling, terrain blocking

- **Layer-by-Layer Spatial Processing**

- Early layers: local station correlations (valley networks)
- Deep layers: regional patterns (Alpine pressure systems)

Result: CRPS 0.64 – modest improvement over Model0, proves attention helps

STGNN2: What the Baseline Does

Architecture: Hybrid TCN-GNN

Approach:

- Stack L layers of: [TCN \rightarrow GNN]

Each layer:

- ① **TCN:** Dilated temporal convolution
 - Captures multi-scale temporal patterns
 - Dilation = 2^l for layer l
- ② **GNN:** Message passing for spatial aggregation
 - Simple GatedGraphNetwork
 - Uniform neighbor weighting
- ③ **Skip:** Store for final multi-scale combination

Final output: sum(all skips) + final layer

STGNN2: Limitations & Solutions

Limitations:

- **Equal Neighbor Weighting**
 - All neighbors treated equally
- **Limited Spatial Reach**
 - Can't capture long-range patterns
 - Requires many hops through graph

Our Solutions:

- **Direction-Aware GATv2**
 - Dynamic neighbor weighting
- **Global Spatial Attention**
 - Direct long-range communication
 - No multi-hop bottleneck

Additional Feature: Smart Ensemble Correction

- Creates two parallel paths:
 - *Shallow*: Simple linear projection
 - *Deep*: Full TCN-GNN-Attention
- Learned gate decides when to use which:

$$\mathbf{h} = g \odot \mathbf{h}_{\text{deep}} + (1 - g) \odot \mathbf{h}_{\text{shallow}}$$

- Trust ensemble when good ($g \approx 0$), correct when biased ($g \approx 1$)

STGNN2: Weather Perspective

How Our Improvements Address Weather Postprocessing:

- **GATv2 Attention**

- Identifies relevant upwind stations for wind correction: upwind stations weighted more heavily
- Captures local effects: valley channeling, terrain blocking

- **Global Attention**

- Long-range dependencies (pressure gradients)
- Example: North Föhn in Ticino driven by Zürich-Lugano pressure difference

- **GRU-inspired Gating**

- Trusts good ensemble predictions
- Prevents over-correction

STGNN2: Results

Method	CRPS ↓	MAE ↓
NWP ensemble	0.98	1.21
Model0 (no graph)	0.67	0.98
STGNN1 (ours)	0.64	0.95
STGNN2 (ours)	0.62	0.91

Key Findings:

- **Best CRPS: 0.62** – significant improvement at *all* lead times
- Outperforms both Model0 and STGNN1
- Consistent gains from short ($t=1h$) to long ($t=96h$) lead times

STGNN3 (TGMM): From ViT to Graph MLP-Mixer Architectural Change

The Journey of Modern Deep Learning:

① Vision Transformers (ViT, 2020)

- Treated images as sequences of **patches** (tokens)
- Used Self-Attention for global context
- **Problem:** $O(N^2)$ complexity

② MLP-Mixer (Google, 2021)

- **Key insight:** Attention isn't necessary!
- Simple MLPs + transposition = global mixing
- **Advantage:** $O(N)$ complexity

③ Graph MLP-Mixer (ICML 2023)

- Adapts "Patch-and-Mix" to **static graphs**
- Uses METIS partitioning to create graph "patches"

④ Temporal Graph MLP-Mixer (2025)

- Extends Graph MLP-Mixer to **dynamic spatiotemporal** data
- Adds **temporal mixing** dimension for time-series

STGNN3 (TGMM): Why MLP-Mixer for Graphs?

Problems with Standard GNNs:

- Over-squashing

- Long-range info compressed into small vectors
- “Traffic jam” for distant signals

- Limited Expressivity

- Bounded by 1-WL test
- Can't distinguish certain graph structures

- Local Receptive Field

- Need L layers for L -hop reach
- Slow information propagation

MLP-Mixer Solution:

- Global Mixing

- Every patch talks to every patch
- “Teleportation” bypasses traffic jam

- High Expressivity

- Exceeds 1-WL test limit
- Distinguishes complex structures

- Linear Complexity

- $O(N)$ vs. $O(N^2)$ for Transformers
- Scalable to large graphs

STGNN3 (TGMM): Creating "Patches" from Graphs

Challenge: Images have grids; graphs have irregular topology.

METIS Partitioning

- Split graph into k clusters
- Minimize inter-cluster edge cuts
- Result: meaningful communities

1-Hop Overlap ("Halo")

- Add 1-hop neighbors to each patch
- Preserves boundary information
- Analogous to image border padding

Weather Station Graphs

- Patches = regional clusters
- Capture regional weather dynamics

Key Insight

Graph patching converts irregular graphs into token sequences, enabling Mixer-style architectures.

STGNN3 (TGMM): The 3D Temporal Mixer

Fundamental Architecture Shift

- **Original T-GMM:** $[B, T_{\text{hist}}, N, F] \rightarrow t+1$ (autoregressive)
- **Our Model:** $[B, T_{\text{lead}}, S, F] \rightarrow \text{all 97h at once}$

3D Mixer Axes

Spatial Mixing	Feature Mixing	Lead-Time Mixing
<ul style="list-style-type: none">• MLP across stations• Regional patterns	<ul style="list-style-type: none">• MLP across variables• Physical interactions	<ul style="list-style-type: none">• MLP across horizon• Temporal evolution

Key Innovation: Parallel Multi-Step Forecasting

Traditional autoregressive approaches predict sequentially (slow, error-prone). Our non-autoregressive method predicts the complete 97h forecast trajectory in **one forward pass**, enabling:

- Faster inference
- No error compounding (each hour predicted independently)
- Consistent trajectories (coherent forecast evolution)

STGNN3 (TGMM): Why Non-Autoregressive

- Weather prediction is an **initial-value problem**
- Model learns $X(t_0) \mapsto \{X(t_0+1h), \dots, X(t_0+97h)\}$ **in one step**
- Conceptually aligned with Numerical Weather Prediction (NWP)

Operational Perspective

Autoregressive Limitations

- Error accumulation
- Inconsistent temporal evolution
- Sequential inference (slow)
- Exposure bias during training

TGMM Advantages

- **Parallel forecasts:** all lead times at once
- **Coherent trajectories:** physically consistent
- **Fast inference:** operationally viable
- **Joint uncertainty:** lead-time covariance

Technical Impact

This architecture enables **operational-grade forecasting** that is fast, consistent, and probabilistic meeting the strict requirements of energy grid management and weather services while maintaining the theoretical advantages of MLP-Mixer efficiency.

STGNN3 (TGMM): Dual-Stream Processing

Original Papers: Single processing stream (patch or node)

Our Adaptation: **Two parallel streams**

Patch Stream

- Input: METIS subgraphs
- View: Coarse-grained / Regional
- Captures: Large-scale weather patterns

Node Stream

- Input: Individual stations
- View: Fine-grained / Local
- Captures: Microclimates
- Example: “Station X is in a valley”

Readout: Combines both streams:

$$\mathbf{h}_{\text{final}} = \text{MLP}([\mathbf{h}_{\text{patch}} \parallel \mathbf{h}_{\text{node}}])$$

Why Both?

Regional context (patches) + local precision (nodes) = best of both worlds

STGNN3 (TGMM): Complete Architecture

Full Pipeline:

- ① **Input:** $\mathbf{X} \in \mathbb{R}^{B \times L \times N \times F}$ + Validity Mask
- ② **METIS Partitioning:** Graph \rightarrow 80–100 patches with 1-hop overlap
- ③ **GNN Encoder:** GINEConv on patches (3 layers, local structure)
- ④ **Dual 3D Mixer:**
 - Patch Mixer (2 layers): Global regional patterns
 - Node Mixer (2 layers): Local station details
- ⑤ **Readout:** Concatenate + MLP $\rightarrow (\mu, \sigma)$ for each station/lead time
- ⑥ **Output:** LogNormal(μ, σ) for 97 lead times \times N stations

Key Properties:

- Linear complexity $O(N)$ – scalable to 1000+ stations
- Global receptive field – captures long-range dependencies
- Probabilistic output

STGNN3 (TGMM): Results

Method	CRPS ↓	MAE ↓
NWP ensemble	0.9808 ± 0.0000	1.2060 ± 0.0000
Model0 (no graph)	0.6715 ± 0.0011	0.9802 ± 0.0037
STGNN1 (ours)	0.6422 ± 0.0046	0.9456 ± 0.0129
STGNN2 (ours)	0.6209 ± 0.0053	0.9077 ± 0.0129
STGNN3 (TGMM)	0.6243 ± 0.0040	0.9017 ± 0.0061

Key Findings:

- Strong MAE performance — highly accurate point predictions
- Competitive CRPS, similar to STGNN2
- Low variance across seeds (± 0.0040)
- Atleast 10x more efficient than comparable architectures

Trade-off

STGNN2: slightly better probabilistic skill (CRPS). TGMM: excellent deterministic accuracy + stability + efficiency.

STGNN3 (TGMM): When to Use?

- Large-scale graphs
- Strict runtime constraints
- Stable and consistent predictions
- Long-range spatial dependencies

The Big Picture

TGMM reflects a broader shift toward **global context with simple MLPs**, avoiding costly attention.

ViT → MLP-Mixer → Graph MLP-Mixer → Temporal Graph MLP-Mixer

Training & Evaluation

Hyperparameter Tuning & Calibration Analysis

Evaluation Protocol

Goal: Compare probabilistic and deterministic performance across models.

Metrics:

- CRPS (probabilistic skill, proper scoring rule)
- MAE (deterministic accuracy)

Aggregation:

- Overall metrics computed across **all stations** and **all lead times**:
- Key lead times: $t = 1, 24, 48, 96$ hours

Robustness:

- 5 independent random seeds
- Report mean \pm std across seeds

Hyperparameter Tuning

Why tuning?

- Deep models are sensitive to hyperparameters
- Manual tuning is inefficient and not reproducible
- We tune **all architectures fairly** under the same protocol

Goal

Automatically select the best configuration for each model using **validation CRPS**, before final multi-seed testing.

Hyperparameter Tuning: How

Framework: Optuna

- **TPE sampler**: intelligent search (not grid / random)
- **MedianPruner**: stops bad trials early
- Short trials (8 epochs) for fast comparison

Why this works

TPE focuses on promising regions of the search space, while MedianPruner avoids wasting time on clearly bad configurations.

Tuning Details & Robustness

Hyperparameters tuned:

- Learning rate: $[10^{-5}, 10^{-3}]$ (log-uniform)
- Hidden channels: $\{32, 64, 128\}$
- Number of layers: $\{1, 2, 3, 4\}$
- Dropout: $[0.1, 0.5]$

Practical considerations:

- 15 successful trials per model
- Automatic batch-size backoff on OOM
- Best config reused for final 5-seed experiments

Final Test Results

Method	CRPS ↓	MAE ↓
NWP ensemble	0.9808 ± 0.0000	1.2060 ± 0.0000
Model0	0.6715 ± 0.0011	0.9802 ± 0.0037
STGNN1 (ours)	0.6422 ± 0.0046	0.9456 ± 0.0129
STGNN2 (default)	0.6648 ± 0.0031	0.9776 ± 0.0080
STGNN2 (ours)	0.6209 ± 0.0053	0.9077 ± 0.0129
STGNN3 (TGMM)	0.6243 ± 0.0040	0.9017 ± 0.0061

Takeaway

Model0 is already strong → carefully designed spatiotemporal graph models yields consistent improvements.

Best overall **CRPS**: STGNN2 (ours). Best overall **MAE**: TGMM.

Lead-Time Breakdown

Why per-lead evaluation matters: forecast uncertainty increases with horizon.

Model	1h	24h	48h	96h
Model0	0.644	0.649	0.668	0.745
STGNN1 (ours)	0.608	0.618	0.638	0.702
STGNN2 (default)	0.628	0.644	0.662	0.724
STGNN2 (ours)	0.578	0.594	0.618	0.687
STGNN3 (TGMM)	0.593	0.603	0.628	0.692

Key observation

All graph-based models improve over the non-graph baseline. Our enhanced STGNN2 achieves the lowest CRPS at all lead times, with the largest gains at short horizons.

Calibration via Rank Histograms

Goal: verify statistical consistency between predicted distributions and observations.




Procedure (per lead time):

- ① For each case: sample **20** values
- ② Insert the observation into the sorted sample set
- ③ Record its rank $\in \{1, \dots, 21\}$ (**21 bins**)


Interpretation:

- Uniform histogram \rightarrow calibrated uncertainty
- U-shape \rightarrow under-dispersed (too narrow)
- Inverted U \rightarrow over-dispersed (too wide)
- Skew \rightarrow bias

Talagrand Diagrams (Lead Time $t = 1\text{h}$)



STGNN2



STGNN3

Why the First Rank is High

Observed pattern: more observations fall below all 20 samples than expected.

Interpretation (intuitive)

When the true wind speed is **very close to zero**, a LogNormal model can struggle to place enough probability mass near zero (it is strictly positive and often right-skewed). So the sampled ensemble values tend to be slightly above the observation, causing the observation to fall into the **lowest rank** more often.

Leads considered: $t = 1, 24, 48, 96$ hours (we show $t = 1\text{h}$ for brevity; longer leads show similar qualitative patterns)

Bonus 1: Graph Refinement (Terrain-Aware kNN)

Motivation: distance-only kNN can create unrealistic links across terrain barriers.

Baseline graph:

- kNN by Haversine distance ($k = 5$)
- Gaussian kernel weights, global θ , threshold = 0.6

Terrain-aware variant:

- Bandwidth from median nearest-neighbor spacing: $\theta_{\text{scale}} = 0.7$
- Sparsity/robustness: $k = 8$, threshold = 0.35, max distance = 120 km, mutual edges
- Orography penalty using **distance to Alpine ridge** to down-weight cross-ridge links

Bonus 1: Results (Graph Ablation on STGNN2)

STGNN2 configuration	CRPS	MAE
Baseline distance graph	0.6209 ± 0.0053	0.9077 ± 0.0129
Terrain-aware graph	0.6173 ± 0.0044	0.9002 ± 0.0105

Takeaway

The gain is small but consistent: terrain-aware connectivity slightly improves both CRPS and MAE, suggesting that reducing unrealistic cross-ridge edges can help performance.